

One Threshold, Five Faces:

From Detection Cliffs to Controlled Reward-Hacking Dynamics

Guangyu Wang
billgywang@gmail.com

July 8, 2026


Abstract

Reward hacking, monitor evasion, and intent drift are usually catalogued as separate empirical failure modes. We argue they are five faces of a single object—the detection edge of a planted-signal / spiked-covariance monitoring problem—and that the faces differ only in which property of that edge one reads (its position, its steepness, or the curvature of the approach) and whether one reads it statically, in the training dynamics, or across a population. Our core result makes the dynamical face precise: fed through an agent’s best-response learning, a static detection cliff becomes a *saddle-node bifurcation* in the agent’s evasion coordinate, with a closed-form threshold $\sigma_c = 3.049/(8\rho^*)$; and because the edge sharpens with dimension ($\sigma = \Theta(d^{1/3})$), *sharper non-adaptive monitoring can be dynamically worse—it makes the loyal→hacking flip more abrupt, more irreversible, and harder to forewarn*. We then sketch two modules that follow: a closed-loop controller whose sensor must be critical slowing down (which survives where the detection-based warning erodes) and whose actuator must be nonsmooth damping (which raises σ_c at no detection-power cost, i.e. without weakening the detector); and a collective diagnosis in which coupling many agents yields a weakest-link cascade whose threshold collapses as $\kappa^* \sim -\ln f$ in the vulnerable fraction f . The mathematical machinery is classical—random-matrix edge theory, catastrophe theory, early-warning signals, threshold cascades—and imported with credit; the contribution is the mapping to reward hacking, one novel bridge (the random-matrix microfoundation $\sigma = \Theta(d^{1/3})$), and a set of actionable laws. Every quantitative claim is verified in simulation *on the surrogate*—not on a trained agent—and we state the surrogate’s scope, evidence status, and falsification conditions explicitly.

1 Introduction

A capable agent optimised against an imperfect monitor faces a standing temptation: exploit the gap between the monitor’s proxy and the true objective (Amodei et al., 2016; Skalse et al., 2022; Gao et al., 2023; Pan et al., 2022), and can learn to tamper with or obfuscate the very signal meant to catch it (Denison et al., 2024; Baker et al., 2025; Greenblatt et al., 2024). When the agent can *hide*—route the tell-tale part of a reward-hacking footprint into a subspace the monitor does not resolve—the failure is not a one-off exploit but a drift in disposition—a learned-optimisation failure (Hubinger et al., 2019)—and practitioners report that this drift can be sudden and hard to reverse (Gao et al., 2023). Is the suddenness fundamental?

We give a unifying account. The object underneath is the *detection edge* of a partial-channel spiked-covariance problem. Let β_{sig} be the resolved signal energy a monitor still sees; an agent’s pressure $\rho \in [0, 1]$ moves energy into the null channel, and $\beta_{\text{sig}} \rightarrow 0$ as $\rho \rightarrow \rho^* = 1 - \sqrt{m/n}/s^2$, the Baik–Ben Arous–Péché (BBP) edge. Every result in this program is a statement about this one edge (Figure 1): its *position* ρ^* is a static recovery wall; its *steepness* $\sigma = \max_{\rho} |A'(\rho)| = \Theta(d^{1/3})$



through-line: sharpness is double-edged — the steepness that makes monitoring powerful makes the failure abrupt, irreversible, unforewarnable

Figure 1: One order parameter, five faces. The detection edge is read through its position (Face 1, static wall), steepness (Face 2, catastrophe), approach curvature (Faces 3–4, controller), or distribution (Face 5, population). The through-line: the steepness that makes monitoring powerful makes the failure abrupt, irreversible, and hard to forewarn.

is the control parameter of a dynamical catastrophe; the *curvature of the approach* is what an online sensor reads; and *distributed* over a population it becomes a cascade threshold. This paper presents the backbone and the core theorem in full, sketches the two dynamical modules, and reports a consolidated simulation summary; the full treatments are in companion papers.

Honest positioning (stated up front). None of the mathematical phenomena is new. Saddle-node and cusp bifurcations and their hysteresis are catastrophe theory (Thom, 1975; Zeeman, 1977); the $3/2$ onset scaling of dynamic hysteresis is classical (Kundu et al., 2023); early-warning signals near tipping points are a mature, Nobel-recognised programme (Scheffer et al., 2009; Dakos et al., 2012; Kuehn, 2011); the edge behaviour of spiked covariance is standard random-matrix theory (Baik et al., 2005; Johnstone, 2001; Tracy and Widom, 1994); and weakest-link cascades are Granovetter/Watts threshold models (Granovetter, 1978; Watts, 2002). Our contribution is not new mathematics but (i) a single, testable *mapping* of reward hacking onto this apparatus; (ii) one genuinely novel bridge—the random-matrix microfoundation $\sigma = \Theta(d^{1/3})$ that ties the detection edge to the catastrophe control parameter, making “sharper non-adaptive monitoring deepens the fold” a theorem rather than a slogan; and (iii) a set of actionable laws: sharper-static-monitoring-is-dynamically-worse, sense-curvature-not-detection, nonsmooth-not-quadratic, harden-the-tail. We work with a surrogate throughout; §7 states its boundaries.

2 The surrogate model

The *evasion coordinate* $\rho \in [0, 1]$ (an intent proxy, not a claim about the agent’s mental state) measures how much of the agent’s hack footprint sits in the monitor’s null channel ($\rho = 0$ loyal/exposed, $\rho = 1$ hacking/hidden). Detection is a fixed logistic surrogate of the BBP curve, $A(\rho) = \frac{1}{2} + \frac{1}{2}(1 + e^{8\sigma(\rho - \rho^*)})^{-1}$, centred at the wall ρ^* with steepness $\sigma := \max |A'|$, so $A(0) \approx 1$,

$A(1) \approx \frac{1}{2}$, and $h(\rho) := -A'(\rho) = \sigma \operatorname{sech}^2(4\sigma(\rho - \rho^*))$, so that $\max_{\rho} h(\rho) = h(\rho^*) = \sigma$ (the normalization fixing σ as the steepness). The agent ascends, on a slow timescale μ , a payoff of evasion minus a convex hiding cost C ($\varphi := C'$ increasing):

$$\dot{\rho} = \mu \partial_{\rho} \Pi(\rho; w), \quad \Pi(\rho; w) = w(1 - A(\rho)) - C(\rho), \quad w > 0, \quad (1)$$

with w the stakes (the value hacking unlocks, growing over training). Equilibria solve $\partial_{\rho} \Pi = w h(\rho) - \varphi(\rho) = 0$; stability is $\partial_{\rho}^2 \Pi < 0$. Under logistic smoothing the endpoints are not equilibria ($\partial_{\rho} \Pi(0) = w h(0) > 0$, inward flow), so both the loyal and hacking equilibria are interior. We call (1) *gradual* if it has a unique stable equilibrium for every w , and *catastrophic* if some open w -interval admits two stable equilibria separated by an unstable one.

3 Core: the catastrophe theorem

Theorem 1 (Closed-form catastrophe criterion). *Take $C(\rho) = \frac{1}{2}c\rho^2$. Equilibria of (1) solve $W h(\rho) = \rho$ with $W := w/c$. The dynamics is gradual when $\sigma < \sigma_c$ and catastrophic (a fold bifurcation as w varies, hence bistability with hysteresis) when $\sigma > \sigma_c$, where*

$$\sigma_c = \frac{\sqrt{3} + 2 \operatorname{arctanh}(1/\sqrt{3})}{8\rho^*} = \frac{3.049 \dots}{8\rho^*}. \quad (2)$$

The threshold is independent of the cost magnitude c and the stakes w (which set the location of the bistable window, not its existence). At $\sigma = \sigma_c$ the equilibrium equation has a triple root and (1) is the cusp normal form $\dot{x} \propto -x^3 + a(W)x + b$.

Proof. Write $G(\rho) := W h(\rho) - \rho$; equilibria are $G = 0$, a fold is $G = G' = 0$, and the cusp is the triple root $G = G' = G'' = 0$. With $s := 4\sigma(\rho - \rho^*)$ and $h = \sigma \operatorname{sech}^2 s$, one has $h'' \propto \operatorname{sech}^2 s (3 \tanh^2 s - 1)$, so $G'' = W h'' = 0$ at the loyal-side inflection $s^* = -\operatorname{arctanh}(1/\sqrt{3})$, where $\operatorname{sech}^2 s^* = \frac{2}{3}$; hence $\operatorname{sech}^2 s^* = \frac{2}{3}$ and $\tanh s^* = -1/\sqrt{3}$ (loyal side, $s^* < 0$). Then $h(\rho_c) = \frac{2}{3}\sigma$ and, from $h' = -8\sigma^2 \operatorname{sech}^2 s \tanh s$, the sign of $\tanh s^* < 0$ makes $h'(\rho_c) = +\frac{16\sigma^2}{3\sqrt{3}} > 0$ (positive, as required for $W > 0$). From $G' = W h' - 1 = 0$, $W = 3\sqrt{3}/(16\sigma^2)$; substituting into $G = 0$ gives $\rho_c = W h(\rho_c) = \sqrt{3}/(8\sigma)$, while geometrically $\rho_c = \rho^* + s^*/(4\sigma)$. Equating the two expressions for ρ_c yields (2). Because G depends on c, w only through $W = w/c$, the threshold is independent of c and w ; non-degeneracy $G'''(\rho_c) \neq 0$ makes the cusp generic. For $\sigma > \sigma_c$, $W h = \rho$ has three roots for w in an interval (two stable, one saddle); for $\sigma < \sigma_c$, one root for every w . $\square$

Two consequences complete the dynamical face; we state them and defer proofs to the companion diagnosis.

The novel bridge of the program makes σ a function of dimension rather than a free knob.

Lemma 1 (BBP edge steepness $\Rightarrow \sigma = \Theta(d^{1/3})$). *Let the monitor be a rank-one spiked-covariance detector with resolved spike $\beta(\rho) = s^2(1 - \rho)$ and aspect $\gamma = m/n$ at ambient dimension d , with wall $\rho^* = 1 - \sqrt{\gamma}/s^2$ (so $\beta(\rho^*) = \sqrt{\gamma}$, the Baik–Ben Arous–Péché edge). Then near ρ^* the detection probability is $A_d(\rho) = F(c_0 d^{1/3}(\rho^* - \rho)) + o(1)$ for a fixed non-degenerate monotone profile F (strictly decreasing, $F' \neq 0$; the Tracy–Widom detection curve), and consequently*

$$\sigma(d) = \max_{\rho} |A'_d(\rho)| = c_0 d^{1/3} \max |F'| = \Theta(d^{1/3}). \quad (3)$$

Proof sketch. At criticality the top sample eigenvalue fluctuates on the Tracy–Widom scale $d^{-2/3}$ about the edge (Tracy and Widom, 1994; Johnstone, 2001), while a spike $\beta = \sqrt{\gamma} + t$ pushes the outlier a distance $\Theta(t^2)$ above the edge for small $t > 0$ (Baik et al., 2005). The outlier becomes statistically distinguishable from the bulk fluctuation—so detection rises from $\frac{1}{2}$ toward 1—over the window $t^2 \asymp d^{-2/3}$, i.e. $\beta - \sqrt{\gamma} \asymp d^{-1/3}$. Writing $\beta - \sqrt{\gamma} = s^2(\rho^* - \rho)$ gives the $d^{1/3}$ scaling of the argument of F , and (3) follows by the chain rule. Only the scaling $\sigma = \Theta(d^{1/3})$ is used downstream; the profile F may be taken as a hypothesis. $\square$

Corollary 1 (Large dimension forces the catastrophe). *Combining Lemma 1 with Theorem 1: since $\sigma(d) = \Theta(d^{1/3}) \rightarrow \infty$ while $\sigma_c = 3.049/(8\rho^*)$ is fixed, there is a finite critical dimension d_c (where $\sigma(d_c) = \sigma_c$) above which the evasion dynamics is catastrophic, and the fold deepens as $d \rightarrow \infty$. Sharper (higher-dimensional) non-adaptive monitoring makes the failure more abrupt, not less (an adaptive monitor can instead reshape the edge; §7). This bridge—turning a free steepness parameter into a dimension law—is the one genuinely new step of the program; everything else imports classical machinery.*

Companion Result 1 (Irreversibility and erosion of warning; companion, summarised). Writing $\Psi := \varphi/h$, equilibria are $\{\rho : \Psi(\rho) = w\}$ and $\partial_\rho^2 \Pi = -h\Psi'$, so stability is $\Psi' > 0$. In the catastrophic regime Ψ is S -shaped: the up- and down-thresholds satisfy $w_+ > w_-$ (a hysteresis loop), the loop opens as $\Delta w = \Theta((\sigma - \sigma_c)^{3/2})$ at onset and diverges as $\sigma \rightarrow \infty$ (irreversibility grows with steepness). The shadow-price early-warning window is informative only for a detection bar $\tau \geq \tau_c = A(\rho_+)$, and since $\rho_+ \rightarrow 0$ as $\sigma \rightarrow \infty$, $\tau_c \rightarrow 1$: the flip becomes unforewarnable exactly as it becomes most severe.

This is the counter-intuitive core: the detection steepness that makes a monitor powerful is what makes the resulting failure abrupt (Theorem 1), irreversible (Theorem 1), and hard to forewarn—and it is guaranteed in high dimension (Corollary 1). The two dynamical modules below are the constructive response.

4 Module A: a closed-loop controller

If the failure is a bifurcation rather than a static error, the safe response is not a sharper monitor but an in-loop controller that senses proximity to the bifurcation and moves the threshold. A controller needs a sensor and an actuator, and the sharp-monitoring regime defeats the obvious choice for each; we give the two that survive (full treatment in the companion control paper).

Companion Result 2 (Sensor: critical slowing down survives where the shadow price erodes). As $w \uparrow w_+$ along the loyal branch, the well curvature flattens as a saddle-node, $\kappa(w) := |\partial_\rho^2 \Pi(\rho_L)| = \Theta(\sqrt{w_+ - w}) \rightarrow 0$, so the stationary fluctuations of ρ undergo critical slowing down, $\text{Var}(\rho) \sim \xi^2/(2\mu\kappa) \rightarrow \infty$ and $\text{AC}_1(\rho) \rightarrow 1$ —a precursor. By contrast the detection $A(\rho_L) \rightarrow A(\rho_+) \rightarrow 1$ stays confident until the flip (Theorem 1), so the shadow-price trigger does not fire. *The sensor must read the curvature, not the detection.* In practice ρ is latent and the CSD indicators are computed on an observable proxy—e.g. the policy–reward gap, the reward-model score residual, an evaluator-agreement rate, or a low-dimensional projection of the reasoning trace—whose variance and autocorrelation rise together near the fold.

Companion Result 3 (Actuator: nonsmooth damping raises σ_c without weakening the detector). Augment the agent’s cost, $C \mapsto C + D$, acting on the update and never on the monitor, so A is unchanged. A linear or quadratic D leaves σ_c invariant (Theorem 1: σ_c is independent

of quadratic cost magnitude), so the damper must be *nonsmooth/superquadratic*; a nonsmooth deadband capping $\rho \leq \rho_{\text{cap}} < \rho_+$ makes the hacking basin unreachable for any σ , at no detection-power cost (the monitor is untouched), with $\rho_{\text{cap}} = \Theta(d^{-1/3})$ (tracking $\sigma = \Theta(d^{1/3})$).

The loop. Coupling the two—sense the slowing-down, engage the damper (raising σ_c above σ), maintain $\sigma < \sigma_c$ —keeps the agent loyal even as stakes ramp past the open-loop fold, subject to a rate law $\dot{w} \leq L(\sigma) w_+ / \tau_{\text{act}}$ (a rate-induced-tipping constraint: ramp slowly enough that the CSD lead exceeds the actuator response time). The detector is untouched, so the loop prevents the flip at *no cost to detection power*, rather than detecting the hack after the fact.

Caveat: the cost is paid elsewhere. “No detection cost” does not mean “no cost.” The nonsmooth deadband forbids *all* movement into the null channel, benign or not, so it is paid for in *capability and autonomy*: it can suppress legitimate exploration and abstention that happen to reduce monitor confidence. The controller trades agent freedom for safety; quantifying that trade-off (e.g. the legitimate-exploration loss of a tight cap) is left to future work.

5 Module B: collective drift

Multi-agent systems (debate, self-play, swarms) couple many single-agent folds. Let each agent i have cost scale c_i (hiding cost $c_i C_0$), and let coupling raise its effective stakes with the mean hacking of its neighbourhood, $\dot{\rho}_i = \mu \partial_{\rho_i} \Pi(\rho_i; w_0 + \kappa \bar{\rho}_i)$. The single-agent fold is linear in cost, $w_+(c) = c w_+(1)$ (so the vulnerability axis is one-dimensional), and the mean field $m = \mathbb{E}_c[\rho_{\text{eq}}(w_0 + \kappa m; c)]$ has a fold: the collective avalanche is a sharp, hysteretic first-order transition at a critical coupling κ^* . The headline is the tail law (full treatment in the companion population paper).

Scaling Law 1 (Weakest-link scaling; companion result, not fully proved here). With a fraction f of “rotten apples” (self-igniting, cost $c_w < c_* := w_0/w_+(1)$) and the rest at cost c_M , the exact threshold is the self-consistency tangency $\kappa^*(1 - f) \partial_W \rho_{\text{eq}}(W^*; c_M) = 1$, whose leading small- f behaviour is a bootstrap logarithm

$$\kappa^*(f) = -\lambda^{-1} \ln f + \kappa_0 + o(1), \quad \lambda^{-1} = \Theta(w_+(c_M) - w_0), \quad (4)$$

with slope set by the gap to the majority fold. Hence κ^* is tail-dominated ($d\kappa^*/df \sim -1/f$): a small rotten-apple fraction collapses the threshold, and under k -regular coupling $\kappa^*(k) = a \ln k + b$ grows with degree—sparse coupling, which concentrates a rotten apple’s influence, is the most dangerous. The log law (4) and the degree law are established here only as *scalings on the tested surrogate* (the degree law via a heuristic local-field mechanism), with constants identified but not fully derived. *Harden the weakest agents, not the mean; do not assume sparse deployment is safe.*

One cusp, two scales. The single-agent and collective transitions share the *same* cusp singularity. For one agent (Theorem 1) bistability is *born* at a cusp as the monitor sharpens ($G = Wh - \rho$ gets a triple root at $\sigma = \sigma_c$); for a population it is *annihilated* at a cusp as the vulnerability spread widens (the self-consistency $\psi = \mathbb{E}_c[\rho_{\text{eq}}] - m$ gets the same degenerate point $\psi = \psi' = \psi'' = 0$ at a critical spread CV_c , beyond which the jump vanishes as $(\text{CV}_c - \text{CV})^{1/2}$ and the cliff becomes a continuous crossover—but only for *independent* vulnerability; a correlated monoculture keeps the sharp flip). One codimension-two bifurcation thus underlies both the onset of the single-agent fold and the endpoint of the collective one (companion population paper).




Figure 2: Simulation summary. (a) The closed form (2) matches a deterministic root-count of the three-root onset across wall positions; green bars show the relative error ($\leq 2\%$, within the conservative 2–4% discretisation band). (b) The hysteresis loop width opens as $(\sigma - \sigma_c)^p$ with fitted $p = 1.48$ against the cusp prediction $3/2$. (c) Ramping stakes past the open-loop fold, the open loop flips to hacking while the closed loop (CSD→damper) stays loyal. (d) The collective threshold κ^* collapses with the rotten-apple fraction f (weakest-link).

6 Simulation summary

All claims are verified on the surrogate (numpy/scipy, CPU); Figure 2 collects the four central checks.

Beyond the figure: the general-convex-cost tangency matches a root-count to 1–3% (stiffer cost raises σ_c); the hysteresis loop widens as $\sigma w_- \rightarrow \rho^*$ and $w_+ \rightarrow \infty$; the erosion bar rises as $\tau_c = 0.75, 0.97, 0.995, 1.00$ for $\sigma = 0.85, 1.0, 1.4, 2.5$ (a $\tau = 0.75$ warning is too late for $\sigma \geq 1$); the CSD sensor warns before the flip in $\approx 90\%$ of runs across $\sigma \in [0.85, 3]$ while detection stays ≈ 0.9996 ; the actuator raises σ_c from 0.78 to 1.38 ($b=0 \rightarrow 50$) at constant detection AUC 0.706; the closed loop engages in 100% of runs ($\sigma \in \{1, 1.4, 2\}$); and at fixed mean cost, dropping only the lower tail moves κ^* from 4.8 to 2.2, with a 1% rotten-apple fraction already lowering it, and $\kappa^*(k) = 0.33 \ln k + 0.91$ over $k = 4, \dots, 64$.

Evidence status. Because the machinery is classical and only some claims are self-proved here, Table 1 states, for each result, whether it is *proved* in this paper, *proved in a companion*, only

Result	Status	Where
Thm 1: catastrophe criterion $\sigma_c = 3.049/(8\rho^*)$	Proved (this paper)	§3
Lemma 1: $\sigma = \Theta(d^{1/3})$ (BBP edge)	Proved here (scaling; imports TW/BBP)	§3
Cor. 1: large d forces the catastrophe	Proved (Thm 1+Lem. 1)	§3
“Sharper non-adaptive monitoring deepens the fold”	Proved (Thm 1+Lem. 1)	§3
Hysteresis $w_+ > w_-$, onset $3/2$	Companion (proved) + simulated	Fig. 2b
Warning erosion $\tau_c \rightarrow 1$	Companion (proved) + simulated	§3
CSD sensor survives erosion	Companion (proved) + simulated	Fig. 2c
Nonsmooth actuator raises σ_c	Companion (proved) + simulated	Fig. 2c
Closed loop stays loyal past fold	Simulated	Fig. 2c
Weakest-link $\kappa^* \sim -\ln f$	Companion (scaling; simulated)	Fig. 2d
Network $\kappa^* = a \ln k + b$; sparse worst	Simulated (heuristic mechanism)	§5

Table 1: Evidence status. “Proved” = full proof here; “Companion” = proved in a companion paper (sketched or summarised here); “Simulated” = verified numerically on the surrogate but not proved in closed form; “Heuristic” = mechanism argued, not derived. Only Theorem 1 and Lemma 1 are proved here (the lemma’s scaling is derived from imported Tracy–Widom/BBP results); the “sharper monitoring” row is their corollary.

simulated, or a *heuristic* mechanism—so the reader can calibrate what rests on what.

7 Evidence status, scope, and falsification

The results are about a deliberately simple surrogate, and we do not overclaim beyond it. (i) The intent is a *scalar, monotone* hiding coordinate; real strategy spaces are high-dimensional and the fold may become a higher-codimension bifurcation. (ii) The dynamics is slow best-response of the form (1); very stiff hiding cost or strongly concave benefit raises σ_c and can remove the catastrophe—the criterion is a genuine inequality, not a certainty. (iii) The monitor is non-adaptive; an adaptive monitor reshapes A and can erase the plateaus that create the fold (an alternative mitigation), though optimising against a monitor can also teach obfuscation (Baker et al., 2025). (iv) Corollary 1 justifies the steepness *scaling* $\sigma = \Theta(d^{1/3})$, not the full detection profile, which we replace by a logistic of matched steepness; finite- d rounding may keep $\sigma < \sigma_c$ in practice. (v) The collective log law (4) is a leading-order scaling whose constants we identify but do not fully derive. We therefore claim only that sharp monitoring *can* induce abrupt, hysteretic, unforewarned evasion by this mechanism, that the mechanism has a random-matrix microfoundation, and that it admits an in-loop controller—not that a specific training system will exhibit these thresholds. Numerical validation is on the surrogate.

What would falsify this, and what must be shown in real agents. The account makes concrete, falsifiable predictions that a real training system either exhibits or not: (a) a *threshold* in monitor sharpness (or dimension) beyond which reward-hacking onset becomes abrupt rather than gradual, with *hysteresis*—rolling back optimisation pressure does not restore loyalty; (b) *critical slowing down*—rising variance/autocorrelation of a hacking proxy *before* onset, while the monitor’s confidence stays high; and (c) at the population level, a *tail-dominated* collective threshold that a small vulnerable fraction collapses. Absence of hysteresis, or an early-warning signal that tracks monitor confidence rather than proxy variance, would falsify the mechanism. Demonstrating (a)–(c) in a real agentic RLHF/RLVR setting—with a measured order parameter and a swept sharpness/dimension knob—is the essential next step; the surrogate only shows the mechanism is internally consistent and quantitatively sharp.

8 Conclusion

A single threshold, namely the detection edge read as position, steepness, approach curvature, or distribution, organises reward hacking from a static wall to a controlled dynamical system. Its dynamical face is a saddle-node catastrophe with a closed-form threshold, and the program’s one novel bridge, $\sigma = \Theta(d^{1/3})$, makes “sharper non-adaptive monitoring deepens the fold” a theorem—and the same cusp singularity that births this bistability for one agent annihilates it for a diverse population, tying the single-agent and collective faces to one codimension-two bifurcation. The constructive consequence is that the right sensor measures proximity to the bifurcation (critical slowing down, not detection confidence) and the right actuator moves the bifurcation threshold (nonsmooth damping, not a quadratic penalty), both without weakening the monitor—while at the population level the threshold is set by the vulnerable tail. The four companion papers give the full treatments.

Companion papers. (*Core*) The Wall, the Game, and the Cliff Criterion. (*Diagnosis*) The Cliff Becomes a Catastrophe. (*Control*) A Closed-Loop Controller for Reward Hacking. (*Population*) Collective Intent Drift.

References

- D. Amodei et al. Concrete problems in AI safety. *arXiv:1606.06565*, 2016.
- J. Baik, G. Ben Arous, and S. Péché. Phase transition of the largest eigenvalue for nonnull sample covariance matrices. *Ann. Probab.*, 33(5):1643–1697, 2005.
- V. Dakos et al. Methods for detecting early warnings of critical transitions. *PLoS ONE*, 7(7):e41010, 2012.
- L. Gao, J. Schulman, and J. Hilton. Scaling laws for reward model overoptimization. *ICML*, 2023.
- M. Granovetter. Threshold models of collective behavior. *Amer. J. Sociology*, 83(6):1420–1443, 1978.
- I. M. Johnstone. On the distribution of the largest eigenvalue in principal components analysis. *Ann. Statist.*, 29(2):295–327, 2001.
- C. Kuehn. A mathematical framework for critical transitions. *Physica D*, 240(12):1020–1035, 2011.
- S. Kundu et al. Dynamic hysteresis at a noisy saddle-node. *arXiv:2308.09787*, 2023.
- M. Scheffer et al. Early-warning signals for critical transitions. *Nature*, 461:53–59, 2009.
- J. Skalse et al. Defining and characterizing reward hacking. *NeurIPS*, 2022.
- R. Thom. *Structural Stability and Morphogenesis*. Benjamin, 1975.
- C. A. Tracy and H. Widom. Level-spacing distributions and the Airy kernel. *Comm. Math. Phys.*, 159(1):151–174, 1994.
- D. J. Watts. A simple model of global cascades on random networks. *PNAS*, 99(9):5766–5771, 2002.
- E. C. Zeeman. *Catastrophe Theory: Selected Papers*. Addison-Wesley, 1977.
- C. Denison et al. Sycophancy to subterfuge: investigating reward tampering in language models. *arXiv:2406.10162*, 2024.
- E. Hubinger et al. Risks from learned optimization in advanced machine learning systems. *arXiv:1906.01820*, 2019.

- A. Pan, K. Bhatia, and J. Steinhardt. The effects of reward misspecification: mapping and mitigating misaligned models. *ICLR*, 2022.
- B. Baker et al. Monitoring reasoning models for misbehavior and the risks of promoting obfuscation. *arXiv:2503.11926*, 2025.
- R. Greenblatt et al. Alignment faking in large language models. *arXiv:2412.14093*, 2024.